

```

Import required libraries

In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

Regression

Explore Dataset of Portland Housing Price

```

In [2]: data = pd.read_csv('portland_housing_dataset.csv',names=['Size','N_bedrooms','Price'])
data.head()
```

Out[2]:

	Size	N_bedrooms	Price
0	2104	3	399900
1	1600	3	329900
2	2400	3	369000
3	1416	2	232000
4	3000	4	539900

```

In [3]: # Load data
data = pd.read_csv('portland_housing_dataset.csv',names=['Size','N_bedrooms','Price'])

# Step 1 : Randomly initialize W
data = np.array(data.values)
X = data[:,::2]
X = np.c_[np.ones(len(X)),X]

X_copy = X.copy()
Y = data[:,2]/1000
```

```

In [4]: plt.title('Number of Bedrooms Vs House Price ')
plt.scatter(X[:,2],Y)
plt.xlabel('Number of Bedrooms')
plt.xticks(ticks=[1,2,3,4,5])
plt.ylabel('House Price (in 1000 $)')
plt.show()
```



```

In [5]: plt.title('Area of House Vs House Price ')
plt.scatter(X[:,1],Y)
plt.xlabel('House area (in sq feet)')
plt.ylabel('House Price (in 1000 $)')
plt.show()
```



Part 1

Implementing LMS (least mean square) Algorithm

```

In [6]: #feature scaling
# it also protect program from overflow error
mean_area = np.mean(X[:,1])
std_area = np.std(X[:,1])
```

```

mean_bedroom = np.mean(X[:,2])
std_bedroom = np.std(X[:,2])
X[:,1] = (X[:,1] - mean_area)/ (std_area)
X[:,2] = (X[:,2] - mean_bedroom)/ (std_bedroom)
```

```

Out[6]: array([[ 1.          ,  0.13141542, -0.22609337],
               [ 1.          ,  0.50956407, -0.22609337],
               [ 1.          ,  0.50790887, -0.22609337],
               [ 1.          , -0.74367706, -1.5543919 ],
               [ 1.          ,  1.27107075,  1.10220517]])
```

```

In [7]: def predict_Y_linear_regression(Input,W) :
Y_pred = Input@W
return Y_pred

def calculate_Jw(Y_pred,Y):
Jw = np.sum( (Y_pred - Y)**2)/2
return Jw

def gradient_Jw(Input,Y,Y_pred):
if len(np.array(Y_pred-Y).reshape(-1,1)) == 1 :
grad_Jw = Input*(Y_pred - Y)
else :
grad_Jw = ( Input.T@ (Y_pred-Y) )
return grad_Jw

def LMS_algorithm(Input,Y,learning_rate=0.001): # Here LMS is also full batch gradient descent
W_arr = []
epsilon = 1e-7
alpha = learning_rate
cntnr = 0
# Step 1 : Randomly Initialize W
W = np.array([700,-150]).reshape(Input[0].shape).reshape(-1,1)

# Step 2 : Loop till convergence
is_converged= False
while not is_converged :
W_arr.append(W)
cntnr += 1
is_converged = True
W_prev = W.copy().reshape(-1,1)
Y_pred = predict_Y_linear_regression(Input,W_prev)
grad_Jw = gradient_Jw(Input,Y,Y_pred).reshape(-1,1)
W_new = W_prev - alpha*grad_Jw
W = W_new
if np.linalg.norm(grad_Jw) > epsilon :
is_converged = False
print("Total no of iterations : ", cntnr)
W_arr = np.array(W_arr).reshape(-1,2)
return W,W_arr
```

```

In [8]: X = X[:,::2]
Y = Y.reshape(-1,1)
W_optimal,W_arr = LMS_algorithm(X,Y,learning_rate=0.01)
W_optimal
```

Total no of iterations : 43

```

Out[8]: array([[340.41265957],
               [105.76413349]])
```

```

In [9]: # Predict House Prices
Y_pred = np.array( list ( map(int, predict_Y_linear_regression(X,W_optimal))) )
Y_pred
```

```

Out[9]: array([354, 286, 394, 261, 474, 338, 277, 263, 256, 272, 332, 340, 325,
               673, 241, 380, 248, 237, 422, 479, 308, 325, 287, 335, 594,
               267, 411, 367, 426, 318, 205, 345, 493, 314, 264, 237, 358, 638,
               362, 295, 372, 416, 232, 185, 320, 233])
```

```

In [10]: Y
```

```

Out[10]: array([[399.9 ],
               [329.9 ],
               [369. ],
               [232. ],
               [539.9 ],
               [299.9 ],
               [314.9 ],
               [198.999],
               [212. ],
               [242.5 ],
               [239.999],
               [347. ],
               [329.999],
               [699.9 ],
               [259.9 ],
               [449.9 ],
               [299.9 ],
               [199.9 ],
               [499.998],
               [599. ],
               [252.9 ],
               [255. ],
               [242.9 ],
               [259.9 ],
               [573.9 ],
               [249.9 ],
               [464.5 ],
               [469. ],
               [475. ],
               [299.9 ],
               [349.9 ],
               [169.9 ],
               [314.9 ],
               [579.9 ],
               [285.9 ],
               [249.9 ],
               [229.9 ],
               [345. ],
               [549. ],
               [287. ],
               [368.5 ],
               [329.9 ],
               [314. ],
               [299. ],
               [179.9 ],
               [299.9 ],
               [235.5 ]])
```

```

In [11]: plt.figure(figsize=(20,5))
plt.title('House Prices')
plt.xlabel('House Area')
plt.ylabel('No of Bedrooms')
plt.yticks([1,2,3,4,5])
plt.scatter(X_copy[:,1],X_copy[:,2], c=Y_pred,marker='x',s=100)
char = plt.colorbar()
char.set_label('House Prices (in 1000 $)',rotation=90)
plt.show()
```



```
In [12]: plt.figure(figsize=(20,5))
plt.title('House Prices')
plt.xlabel('House Area')
plt.ylabel('Houses Price')
plt.scatter(X_copy[:,1],Y_pred, c=Y_pred,s=100,label='Predicted House Prices')
cbar = plt.colorbar()
cbar.set_label('House Prices (in 1000 $)',rotation=90)
plt.scatter(X_copy[:,1],Y,marker='*',s=100,label='True House Prices')
plt.legend()
plt.show()
```



Part 2

Visualize Contours of J(W) for different values of learning rate η

```
In [13]: # Draw contours of Loss Function for different learning rates
fig, ax = plt.subplots(3,3,figsize=(15,15))
plt.suptitle('Contours of J(W) for different learning rates ')
learning_rates = [0.0005, 0.0008, 0.001, 0.005, 0.005, 0.01 , 0.02, 0.025, 0.03]

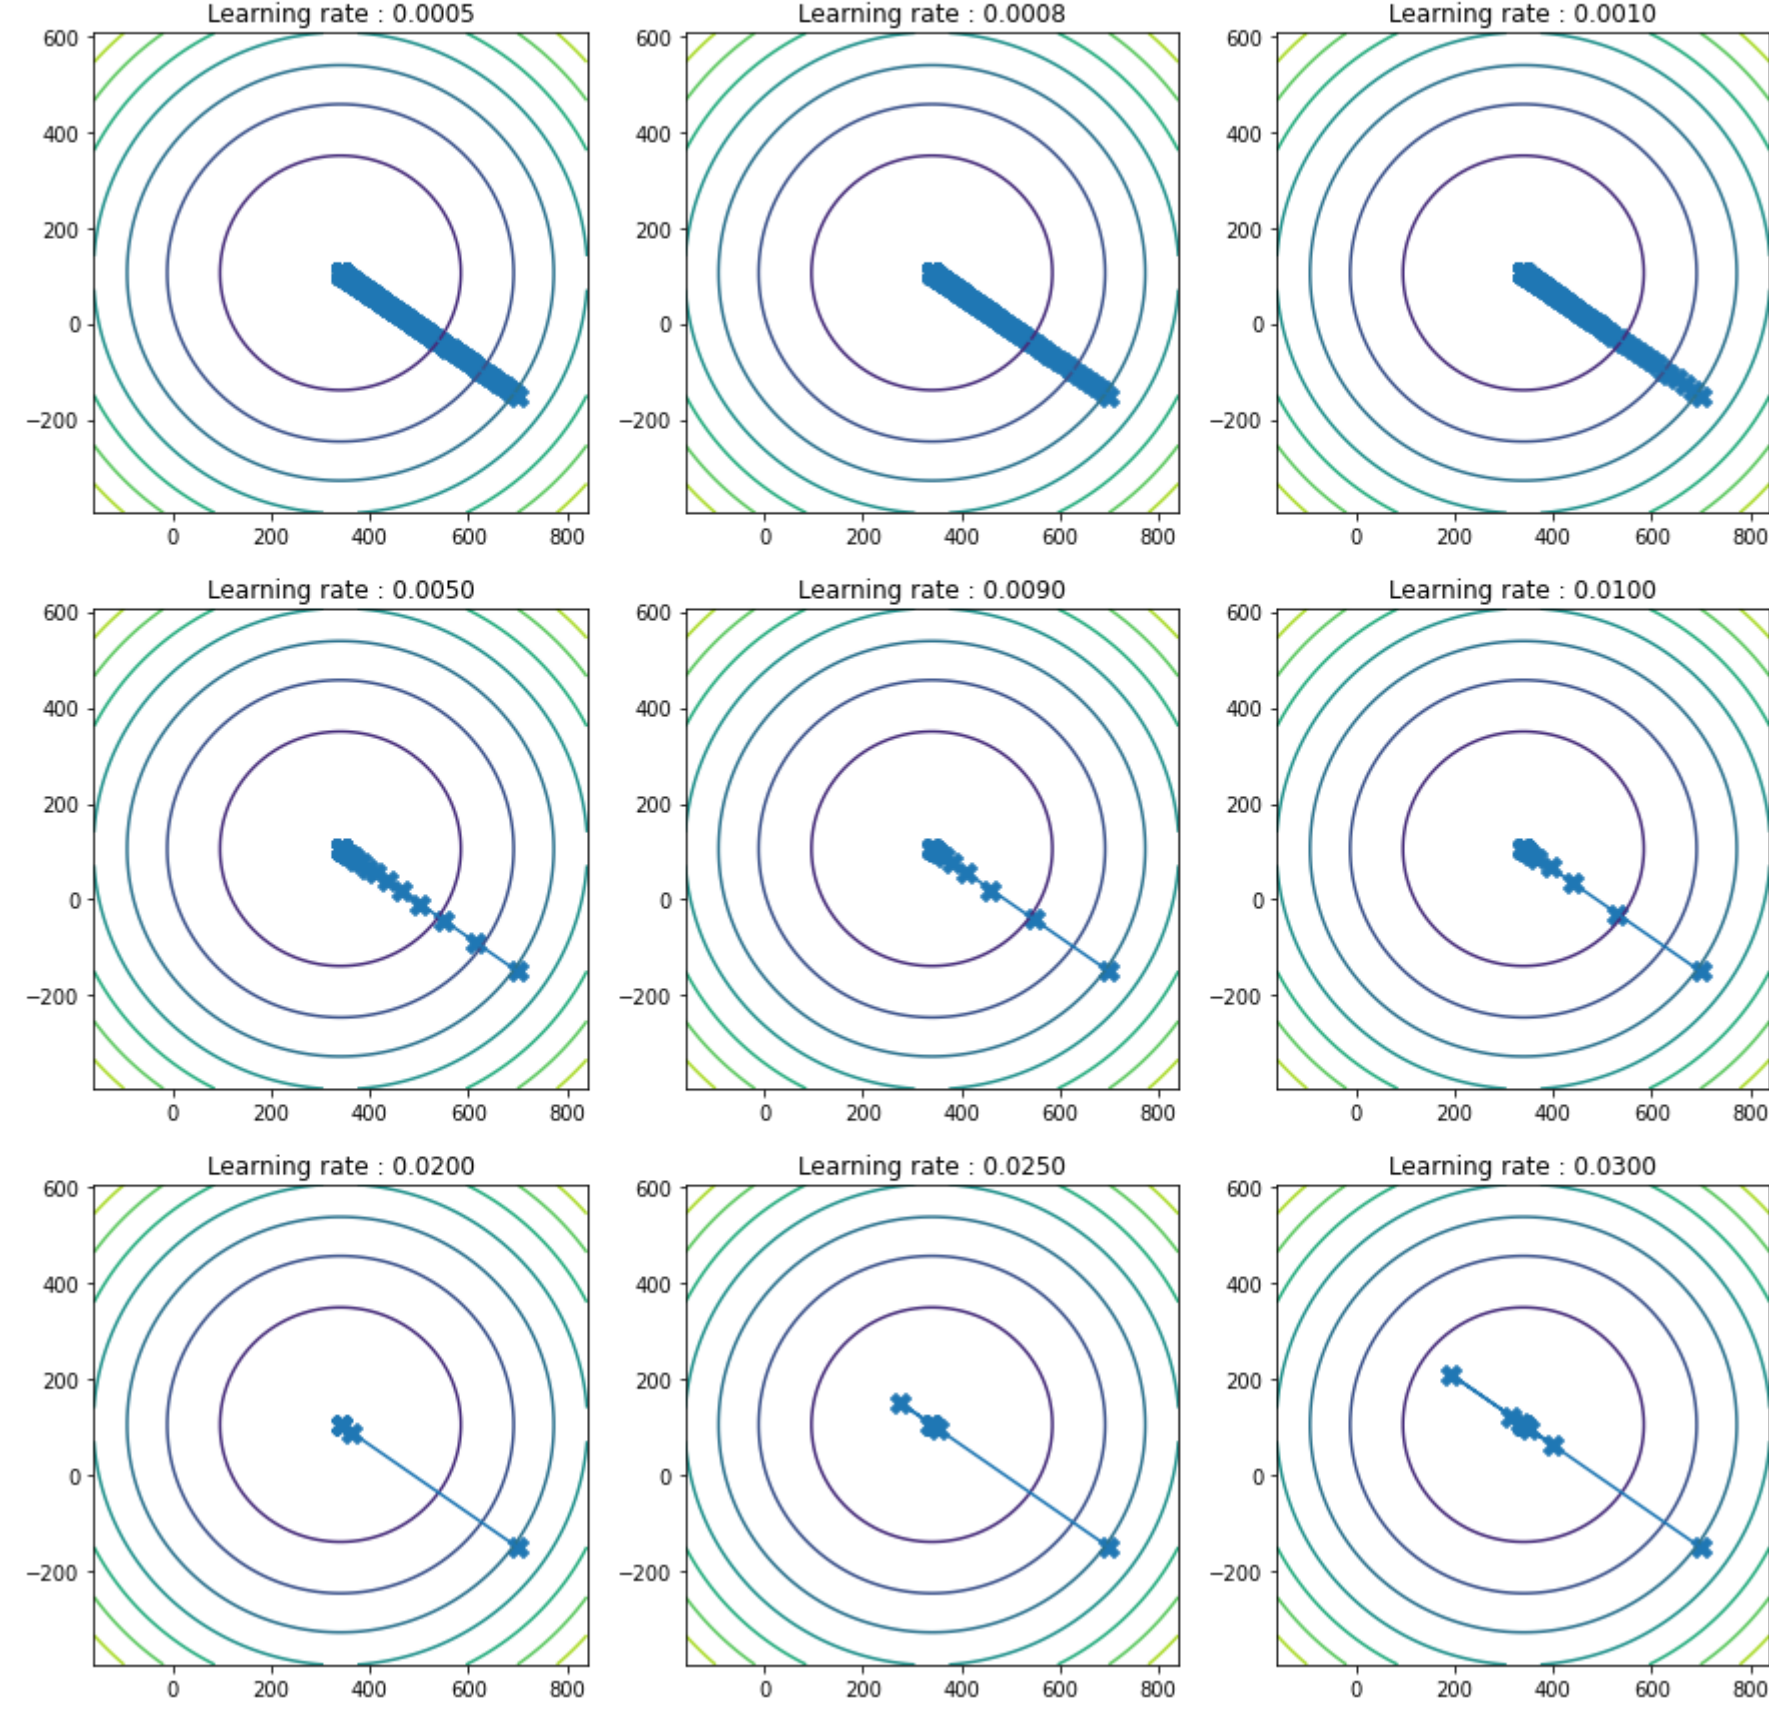
for idx, learning_rate in enumerate(learning_rates) :
    W_optimal,W_arr = LMS_algorithm(X,Y,learning_rate)
    x = np.linspace(W_optimal[0]-500, W_optimal[0]+500, 100)
    y = np.linspace(W_optimal[1]-500, W_optimal[1]+500, 100)
    xx,yy = np.meshgrid(x,y)

    Z = []
    for a in x :
        W_arr = []
        for b in y :
            W = [a,b]
            W = np.array(W).reshape((-1,1))
            Y_pred = X@W
            Jw = calculate_Jw(Y_pred, Y)
            W_arr.append(W)
            Z.append(Jw_arr)
        del W_arr
    Z = np.array(Z)
    i,j = idx//3 , idx%3
    print(i,j)

    ax[i,j].set_title('Learning rate : %0.4f'%learning_rate)
    ax[i,j].contour(xx,yy,Z)
    ax[i,j].scatter(W_arr[:,0],W_arr[:,1],marker='X',s=100)
    ax[i,j].plot(W_arr[:,0],W_arr[:,1])
```

Total no of iterations : 1097
0 0
Total no of iterations : 681
0 1
Total no of iterations : 543
0 2
Total no of iterations : 99
1 0
Total no of iterations : 49
1 1
Total no of iterations : 43
1 2
Total no of iterations : 11
2 0
Total no of iterations : 16
2 1
Total no of iterations : 31
2 2

Contours of J(W) for different learning rates



Part 3

Implementing Stochastic Gradient Descent

```
In [14]: def stochastic_gradient_descent(Input,Y,learning_rate=0.02,batch_size=5): # Here LMS is also full batch gradient descent
    W_arr = []
    epsilon = 1e-2
    alpha = learning_rate
    cntnr = 0
    # Step 1 : Randomly Initialize W
    W_prev = 100000000
    W = np.array([700,-150]).reshape(Input[0].shape)
    # Step 2 : Loop till convergence
    is_converged= False
    while not is_converged :
        cntnr += 1
        is_converged = True
        for i in range(min(len(Input),batch_size)):
            W_prev = W.copy()
            Y_pred = predict_Y_linear_regression(Input[i],W_prev)
            grad_W = gradient_Jw(Input[i],Y[i],Y_pred)
            W_new = W_prev - alpha*grad_W
            W = W_new

        W_arr.append(W)
        Jw_cur = calculate_Jw(predict_Y_linear_regression(Input,W),Y)
        tmp = np.sum(np.abs(Jw_cur - Jw_prev))
        print(tmp)
        if tmp> epsilon :
            is_converged = False
        Jw_prev = Jw_cur
        print("Batch size : %d , learning_rate(lr) : %f"% (batch_size,learning_rate))
        print("Total no of iterations : ", cntnr)
    W_arr = np.array(W_arr).reshape(-1,len(X[0]))

    return W , W_arr
```

```
In [15]: W_optimal,W_arr = stochastic_gradient_descent(X,Y,learning_rate=0.01,batch_size=5)
W_optimal
```

```
Out[15]: array([356.88570766, 130.12654783])
```

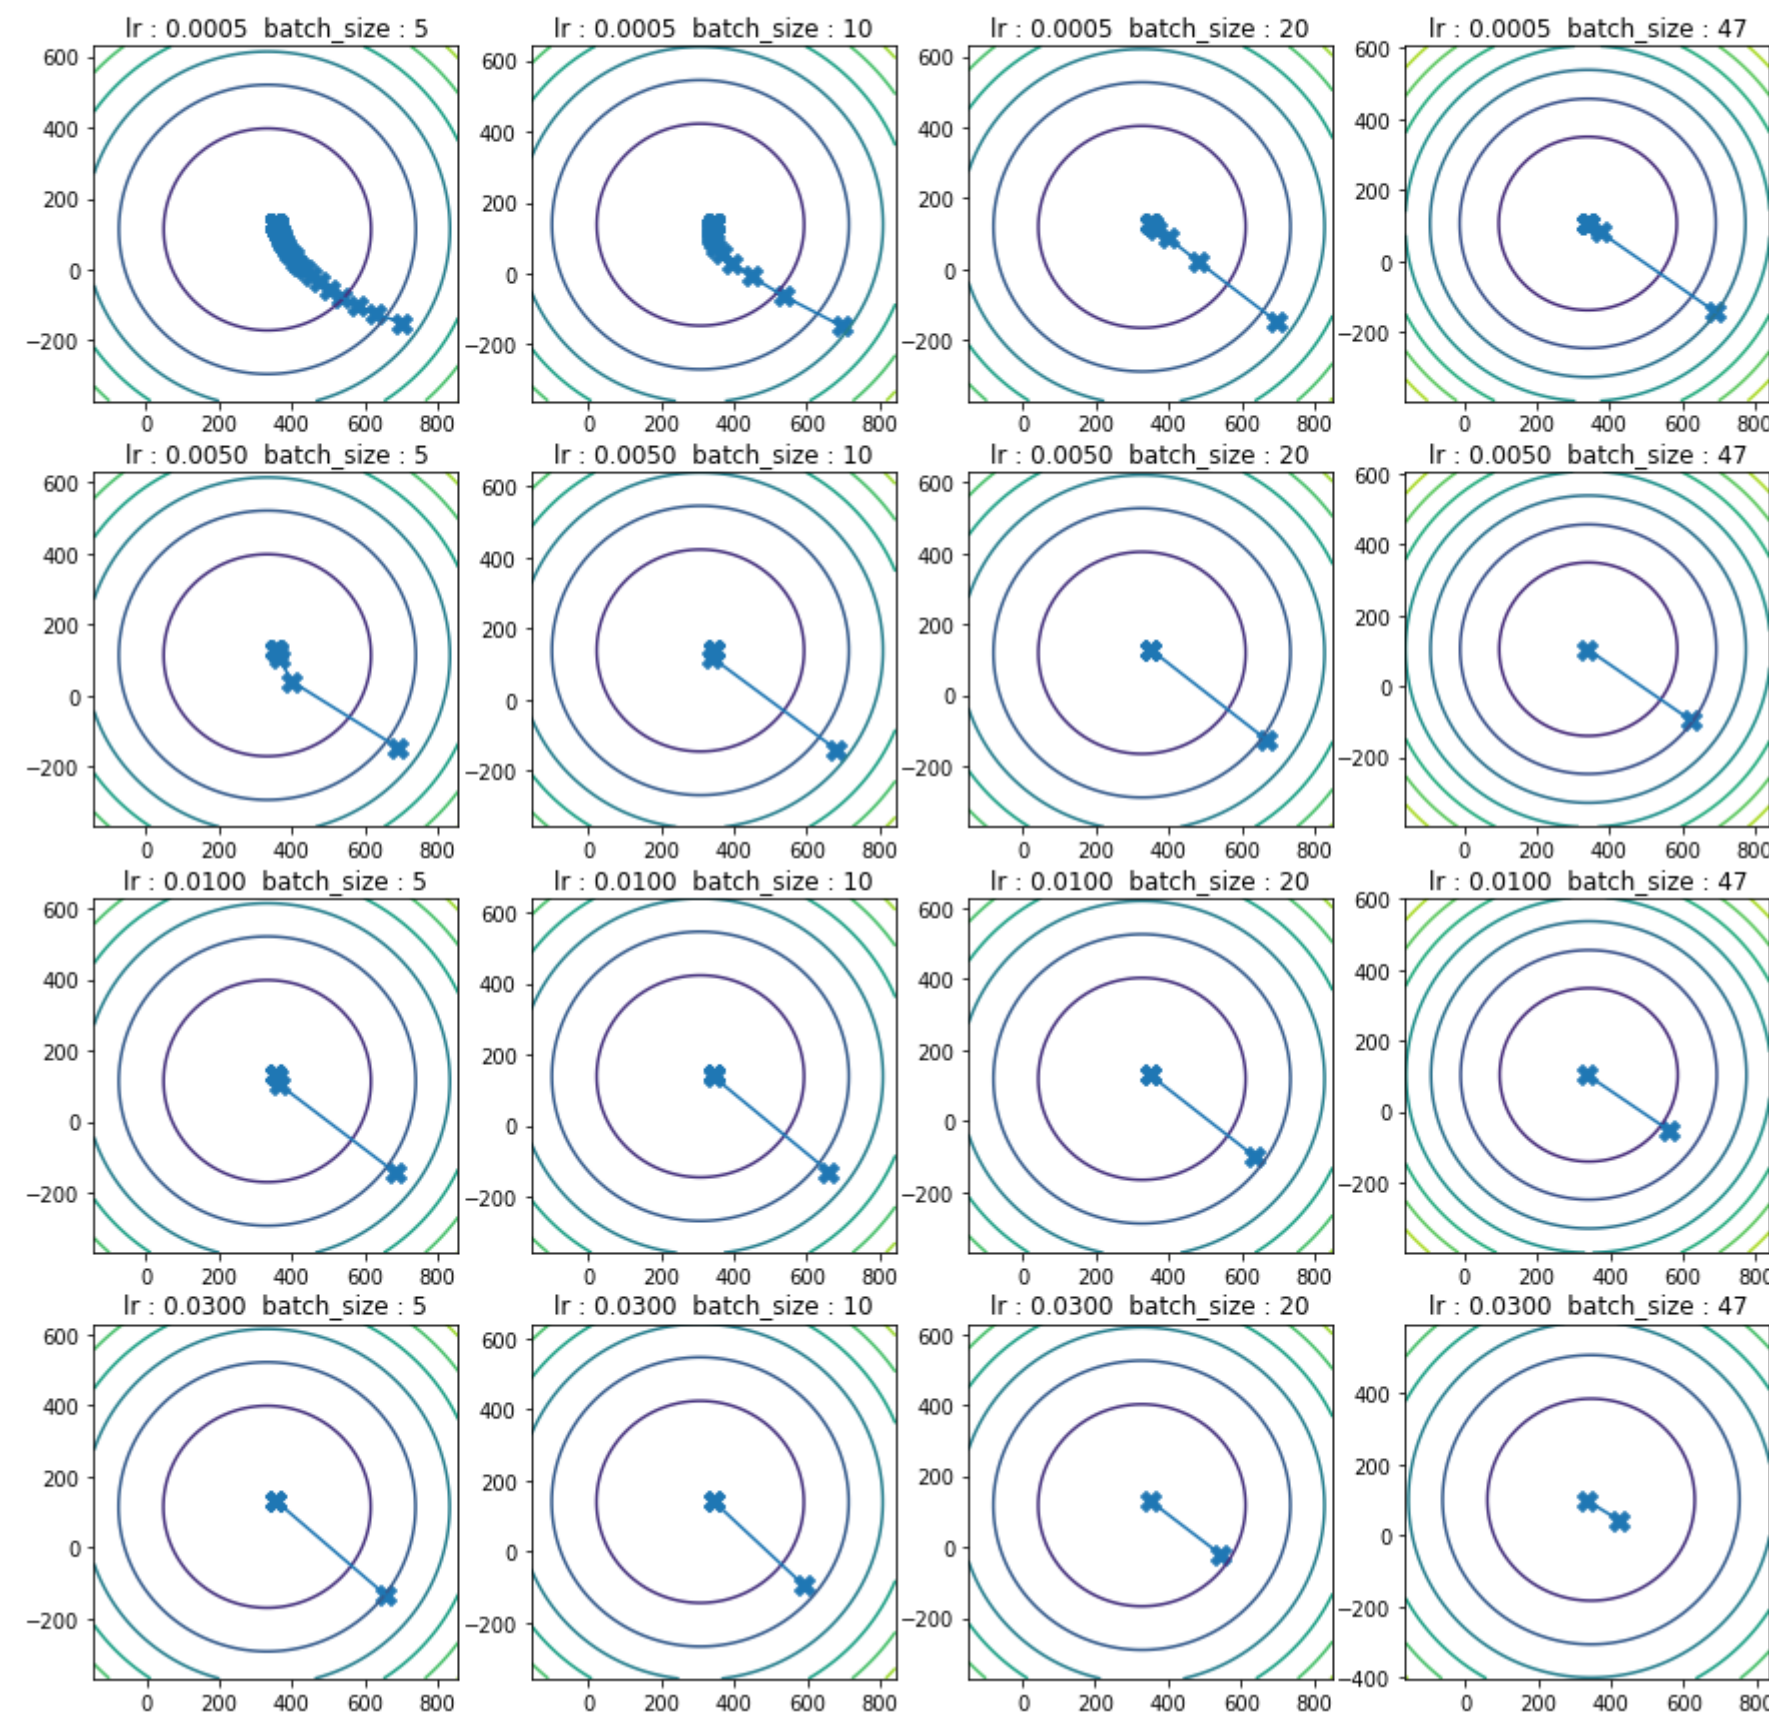

In [16]: `# Draw contours of Loss Function for different learning rates and batch size in stochastic gradient descent`

```
fig, ax = plt.subplots(4,4,figsize=(15,15))
plt.suptitle('Contours of J(W) for different learning rates and Batch Sizes ')
learning_rates = [0.0005, 0.005, 0.01, 0.03]
batch_sizes = [5,10,20,47]
for l,learning_rate in enumerate(learning_rates):
    for j, batch_size in enumerate(batch_sizes):
        W_optimal,W_arr = stochastic_gradient_descent(X,Y,learning_rate,batch_size)
        x = np.linspace(W_optimal[0]-500, W_optimal[0]+500, 100)
        y = np.linspace(W_optimal[1]-500, W_optimal[1]+500, 100)
        xx,yy = np.meshgrid(x,y)
        tmp = [ W_arr[i] for i in range(0,len(W_arr)-1,100)] # pick only 100 intermediate W's to plot progress
        tmp.append(W_arr[-1])
        del W_arr
        W_arr = np.array(tmp)

        Z = []
        for a in x:
            W_arr = []
            for b in y:
                W = np.array(W).reshape(-1,1)
                Y_pred = X@W
                Jw = calculate_Jw(Y_pred, Y)
                Z.append(Jw)
            del W_arr
        Z = np.array(Z)

        ax[l,j].set_title('lr : %0.4f batch_size : %d'%(learning_rate,batch_size))
        ax[l,j].contour(xx,yy,Z)
        ax[l,j].scatter(W_arr[:,0],W_arr[:,1],marker='x',s=100)
        ax[l,j].plot(W_arr[:,0],W_arr[:,1])
```

Contours of J(W) for different learning rates and Batch Sizes



Part 4

Implementing Locally Weighted Linear Regression

```
In [17]: def weight_matrix(X, x, tau = 1):
    return np.exp( -1 * (X[:,1:2]-x)**2) / ((2**tau)**2)

def predict_Y_linear_regression(Input,W):
    Y_pred = Input@W
    return Y_pred

def calculate_LW_Jw(Y_pred,Y,x):
    W = weight_matrix(X,x).reshape(-1,1)
    Jw = np.sum(W*(Y_pred - Y)**2)/2 #Locally weighted cost function
    return Jw

def gradient_LW_Jw(Input,Y,Y_pred,x):
    W = weight_matrix(Input,x)
    # print(W)
    if len(np.array(Y_pred-Y).reshape(-1,1)) == 1:
        grad_Jw = Input*((Y_pred - Y)*W)
    else:
        grad_Jw = ( Input.T@((Y_pred-Y)*W) )
    return grad_Jw

def LW_LMS_algorithm(Input,Y,learning_rate=0.001,x = 1): # Here LMS is also full batch gradient descent
    W_arr = []
    epsilon = 1e-7
    alpha = learning_rate
    cntnr = 0
    # Step 1 : Randomly Initialize W
    W = np.array([700,-150]).reshape(Input[0].shape).reshape(-1,1)

    # Step 2 : Loop till convergence
    is_converged= False
    while not is_converged:
        W_arr.append(W)
        cntnr += 1
        is_converged = True
        W_prev = W.copy().reshape(-1,1)
        Y_pred = predict_Y_linear_regression(Input,W_prev).reshape(-1,1)
        grad_Jw = gradient_LW_Jw(Input,Y,Y_pred,x = x)
        # print(grad_Jw.shape, W_prev.shape, W.shape, Y_pred.shape, Y.shape)
        W_new = W_prev - alpha*grad_Jw
        W = W_new
        # print(np.linalg.norm(grad_Jw))
        # print(grad_Jw)
        if np.linalg.norm(grad_Jw) > epsilon:
            is_converged = False
        print("Total no of iterations : ", cntnr)
        W_arr = np.array(W_arr).reshape(-1,2)
    return W,W_arr

X = X[:,1:2]
x = 1500 # POI
x = (x - mean_area)/std_area # de-scale
print(x)
Y = Y.reshape(-1,1)
W_optimal_lw, _ = LW_LMS_algorithm(X,Y,learning_rate=0.01, x = x)
print(W_optimal_lw)
W_optimal_lms, _ = LMS_algorithm(X,Y,learning_rate=0.01)
print(W_optimal_lms)

# Make predictions of locally weighted
X_ = sorted(X, key = lambda x : x[1])
X_ = np.array(X_)

Y_pred_lw = X_@W_optimal_lw
Y_pred_lms = X_@W_optimal_lms

X_ = std_area*X_ + mean_area

plt.figure(figsize=(20,5))
plt.title('House Prices')
plt.xlabel('House Area')
plt.ylabel('House Price')
plt.scatter(X[:,1],Y,marker='*',s=100,label='True House Prices')
plt.plot(X[:,1],Y_pred_lw,c = 'red', label='Predicted House Prices (Locally Weighted)')
plt.plot(X[:,1],Y_pred_lms, c= 'green', label='Predicted House Prices')

cbar = plt.colorbar()
cbar.set_label('House Prices (in 1000 $)',rotation=90)
plt.legend(bbox_to_anchor=(1.5,0.8))
plt.show()

-0.6368343721168225
Total no of iterations : 239
[[335.73239883]
 [ 98.26023344]]
Total no of iterations : 43
[[348.41265957]
 [185.76413349]]
```

